

A Convolutional Neural Network (CNN) for handwritten digit recognition on FPGA using HLS

Giacomo Boldini

University of Parma
Master's Degree in Computer Science
Embedded Systems

AA 2021/22

Project repository here

[1]

Goals & Outline

Goals

- Creation of a NN for handwritten digit classification.
- Implementation of the NN on FPGA using HLS/Vivado.
- Prove that HW solutions is faster than SW (C) solutions.

Outline:

- ① Python: Create and train NN model
- ② C: NN implementation
- ③ Vitis/C++: NN synthesis and validation
- ④ Conclusions


The (C)NN model

CNN architecture choosed:

- image-processing task;
- no need of manual feature extraction: done automatically;
- less number of parameters than other NNs.

API: Python Keras/Tensorflow [2].

Model as simple as possible:



Other model configurations

Same model but different number of filters:

(epochs = 50, validation split = 0.2)				
filter number	# param	last val. acc.	test acc.	pred. time (ms)
8	15k	97.78	98.07	36~38
16	31.5k	98.14	98.27	38
32	63k	98.17	98.37	38~40
64	126k	98.32	98.38	38~40

Same model but different number of filters + 1 dense layer:

(epochs = 50, validation split = 0.2)				
filter number	# param	last val. acc.	test acc.	pred. time (ms)
8	158k	92.67	93.18	38
16	315k	93.80	93.42	50
32	630k	94.98	95.05	50
64	1256k	94.73	94.23	52

Training

TrainX shape = (60000, 28, 28)

Training epochs = 10 (empiric)

Layers' trainable parameters:



Layer (type)	Output Shape	Param #
ZeroPadding2D	(30, 30, 1)	0
Conv2D	(28, 28, 8)	80
MaxPooling2D	(14, 14, 8)	0
Flatten	(1568)	0
Dense	(10)	15690
TOT		15770

Accuracy:

- validation set (20% of test set): 97.78%
- test set (#10000 samples): **98.070%**

Mean time for a prediction: **~35 ms**


Training history:



C cnn() function

```
void cnn(float img_in [IMG_ROWS][IMG_COLS], float prediction[DIGITS])
{
    // Normalization and padding.
    float pad_img [PAD_IMG_ROWS][PAD_IMG_COLS] = { 0 };
    normalization_and_padding(img_in, pad_img);
    // Convolution.
    float features [FILTERS][IMG_ROWS][IMG_COLS] = { 0 };
    convolutional_layer(pad_img, features);
    // Pooling.
    float pool_features [FILTERS][POOL_IMG_ROWS][POOL_IMG_COLS] = { 0 };
    max_pooling_layer(features, pool_features);
    // Flattening.
    float flat_array [FLAT_SIZE] = { 0 };
    flattening_layer(pool_features, flat_array);
    // Dense.
    dense_layer(flat_array, prediction);
}
```

C main() / testbench



- MNIST TestX samples: 10000 N: 100 ~ 250
- Accuracy: $\frac{\text{correct predictions}}{\text{total predictions}}$ Test successful \Leftrightarrow Accuracy $\geq 95\%$

Mean time for a prediction:

- **0.82 ms** - O0 (~40x faster than Python)
- **0.17 ms** - O3 (~200x faster than Python)

Code optimizations for FPGA

CNN do not need all the data from the previous layer to start computing the output response for the current layer [3].

C implementation not optimized for FPGA deployment:

- data input/output with std array;
- does not create/support parallelism (dataflow).

Optimize code:

- 1 `hls::stream` [4] between functions:
FIFO with blocking API `read()` and `write()`.
- 2 New function `dataflow_section(img1, img2, ..., img8)`:
Clone input image `FILTER_NUMBER` times.

C simulation

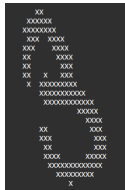
Total predictions: 500.

Correct predictions: 98.20 % → **OK**.

Average latency: 2.33 ms → a little bit more than C.

Some bad classifications:


(images normalized and rounded)



Expected: 3

Got:

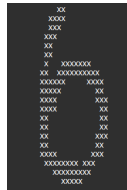
0: 0.000002
1: 0.000000
2: 0.001373
3: 0.213332
4: 0.000003
5: 0.000935
6: 0.000000
7: 0.000000
8: 0.783027
9: 0.001329



Expected: 4

Got:

0: 0.000000
1: 0.000045
2: 0.000020
3: 0.000661
4: 0.253086
5: 0.000059
6: 0.000414
7: 0.000036
8: 0.000321
9: 0.745357



Expected: 6

Got:

0: 0.735325
1: 0.000000
2: 0.000000
3: 0.000000
4: 0.000000
5: 0.000019
6: 0.264633
7: 0.000000
8: 0.000004
9: 0.000020

C synthesis I

Common parameters:

- Target device: **xc7a200tfbg484-1**
- Target clock period: **10ns** (clock freq.: 100 MHz)

Different “*levels of optimization*” (directives):

① No directives

Target	Estimated	Uncertainty
10.00 ns	7.300 ns	2.70 ns

Modules & Loops	Issue Type	Violation Type	Distance	Slack	Latency(cycles)	Latency(ns)	Iteration Latency	Interval	Trip Count	Pipelined	BRAM	DSP	FF	LUT	URAM
▼ cnn				-	361721	3,617E6	-	361722	-	no	231	69	16679	18273	0
> dataflow_section				-	342698	3,427E6	-	342698	-	no	215	69	16409	17588	0
> pad_for_rows				-	17160	1,720E5	572	-	30	no	-	-	-	-	-
> VITIS_LOOP_102_1				-	1860	1,860E4	62	-	30	no	-	-	-	-	-

C synthesis II

2 Default directives

Target	Estimated	Uncertainty
10.00 ns	7,300 ns	2.70 ns

Modules & Loops	Issue Type	Violation Type	Distance	Slack	Latency(cycles)	Latency(ns)	Iteration Latency	Interval	Trip Count	Pipelined	BRAM	DSP	FF	LUT	URAM
▼ cnn				-	19523	1,950E5	-	19524	-	no	264	104	40089	30930	0
> ● cnn_Pipeline_pad_for_rows_pad_for_cols				-	920	9,200E3	-	920	-	no	0	0	1018	419	0
> ● cnn_Pipeline_VITIS_LOOP_102_1_VITIS_LOOP_103_2				-	903	9,030E3	-	903	-	no	0	0	45	167	0
> ● dataflow_section		II Violation		-	17695	1,770E5	-	17695	-	no	248	104	39017	29916	0


3 Dataflow directive

Target	Estimated	Uncertainty
10.00 ns	7,241 ns	2.70 ns


Modules & Loops	Issue Type	Violation Type	Distance	Slack	Latency(cycles)	Latency(ns)	Iteration Latency	Interval	Trip Count	Pipelined	BRAM	DSP	FF	LUT	URAM
▼ cnn				-	5825	5,825E4	-	5826	-	no	384	129	46902	38213	0
> ● cnn_Pipeline_pad_for_rows_pad_for_cols				-	920	9,200E3	-	920	-	no	0	0	1018	419	0
> ● cnn_Pipeline_VITIS_LOOP_102_1_VITIS_LOOP_103_2				-	903	9,030E3	-	903	-	no	0	0	45	167	0
▼ dataflow_section				-	3997	3,997E4	-	3998	-	dataflow	368	129	45828	37234	0
> [X] convolutional_layer				-	3997	3,997E4	-	3998	-	dataflow	0	80	20585	18273	0
> [X] max_pooling_layer				-	792	7,920E3	-	793	-	dataflow	0	0	3513	4049	0
> [X] flattening_layer				-	198	1,980E3	-	199	-	dataflow	0	0	353	825	0
> [X] dense_layer				-	2013	2,013E4	-	2014	-	dataflow	320	40	15584	9632	0
> ● dense_layer_soft_max				-	871	8,710E3	-	871	-	no	0	9	1071	1689	0

C synthesis III

Dataflow view:



(zoom on convolutional_layer)



Validation and implementation

C/RTL Cosimulation → OK

Modules & Loops		Avg II	Max II	Min II	Avg Latency	Max Latency	Min Latency
▼	● cnn	6747	6747	6747	6746	6746	6746
>	● cnn_Pipeline_pad_for_rows_pad_for_cols	6747	6747	6747	918	918	918
>	● cnn_Pipeline_clone_for_rows_clone_for_cols	6747	6747	6747	901	901	901
>	⊗ dataflow_section	6747	6747	6747	4922	4922	4922

Total predictions: 100
 Correct predictions: 99.00 %
 Average latency: 0.290000 (ms)
 *** C/RTL co-simulation finished: PASS ***

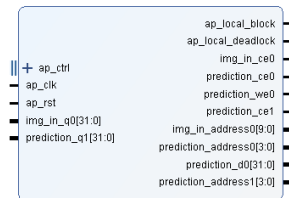
→ prediction time: **0.067 ms**

Implementation (Vivado)

	Verilog
SLICE	12940
LUT	26381
FF	38178
DSP	129
BRAM	224
URAM	0
LATCH	0
SRL	1007
CLB	0


	Verilog
CP required	10.000
CP achieved post-synthesis	8.123
CP achieved post-implementation	9.449

Timing met



Conclusions

Main goal reached: **HW faster than SW.**



But, as future works:

- grid-search on NN architecture could increase accuracy ($>$ performance) and reduce FPGA area ($<$ price);
- using *fixed-point* arithmetic could reduce area;
- small SW changes could improve parallelism.

References

- [1] *Github: HLS-CNN*. [Project repository]. URL: <https://github.com/FedericoSerafini/HLS-CNN>.
- [2] Francois Chollet et al. *Keras*. 2015. URL: <https://github.com/fchollet/keras>.
- [3] Duda S. *How to Implement a Convolutional Neural Network Using High Level Synthesis*. Ed. by amiq.com. [Online; posted 14-December-2018]. 2018. URL: <https://www.amiq.com/consulting/2018/12/14/how-to-implement-a-convolutional-neural-network-using-high-level-synthesis/>.
- [4] *Vitis High-Level Synthesis User Guide: HLS Stream Library*. [Online; visited june-2022]. URL: <https://docs.xilinx.com/r/en-US/ug1399-vitis-hls/HLS-Stream-Library>.

Thank you for your attention.